

Identifying Customer Needs from User-Generated Content

Artem Timoshenko (Northwestern), John R. Hauser (MIT)
Marketing Science, 2019

Presenter: Shuyi Zhang

May 14, 2025





Motivation:

- ▶ Traditional Voice-of-Customer (VOC) methods are time-consuming and costly.
 - Identification: Qualitative Interviews
 - Structuring: Manual reviews & summarize by multiple analysts.
- ▶ User-Generated Content (UGC) is an underutilized resource for product development.
 - Abundant unstructured textual data e.g. reviews, speeches
- ▶ Machine learning (ML) may offer efficiency gains in extracting valuable customer insights.

General Research Question: Can we utilize ML to effectively identify **customer needs** from UGC? If so, how well does it perform?

Model: Summary of the System Architecture



Figure 1. System Architecture for Identifying Customer Needs from UGC

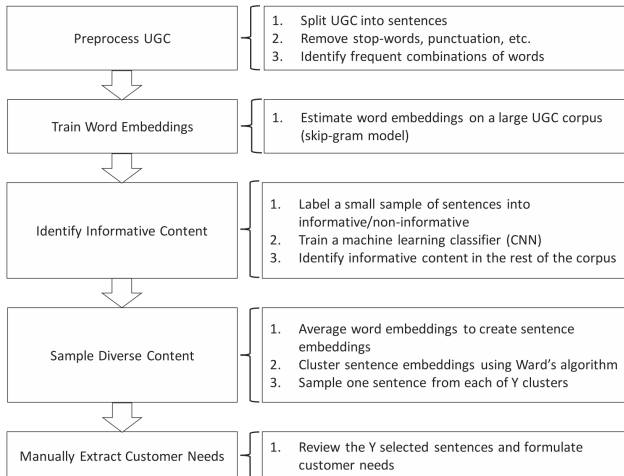


Figure: Timoshenko and Hauser (2019) Figure 1



- ▶ 115,099 oral-care reviews on Amazon spanning the period from 1996 to 2014, randomly sampled 12,000 sentences split into an initial set of **8,000 sentences** and a second set of 4,000 sentences.
- ▶ For the 8000 sentences, hired professional marketing analysts to fully **code every sentence to determine whether it contained a customer need** and, if so, whether the customer need could be mapped to a customer need identified by the VOC, or whether the customer need was a newly identified customer need.

To summarize, the data I focus on for the replication are:

- ▶ `all_sentences`: All the oral care reviews.
- ▶ `8000_sentences`: Fully labeled sentences. `Informative = 1`



0-First of all, Target.com is selling the wrong unit.

1-:-)I'm achieving much better results with the Sonicare than with manual brushing, for 2 reasons: First, there's no way I come close to 31,000 brush strokes per minute by hand.

1-It's constantly getting dirty with dust and tooth paste.

0-I finally got this toothbrush after I have seen alot of people use them.

Model: 1. Preprocessing UGC



Steps:

1. Unsupervised tokenizer (nltk®ex) to eliminate stop-words(e.g. the, and) and symbols
2. Join 'frequently together' words into Phrases (e.g. even_though)
3. "We drop sentences that are less than four words or longer than 14 words after preprocessing."
4. Train Word Embeddings with **Skip-Gram Model**

Model: 1. Preprocess UGC



Let's look at the following example¹ from *8000_sentences.csv*:

Well I really don't know how well the 4-pack of Teledyne piks (BRJ4) works because, even though it's clearly what I ordered, Goodman's, without notice, sent me the Official WaterPik 2-Pack (JT-70E) in its place for the same total cost to me as what I had paid for the Teledyne 4-pack. Two piks in place of four?.

After the cleaning and tokenizing, we have:

► ['well', 'really', 'know', 'well', 'pack', 'teledyne', 'piks', 'brj', 'works', 'even_though', 'clearly', 'ordered', 'goodman', 'without', 'notice', 'sent', 'official', 'waterpik', 'pack', 'jt', 'e', 'place', 'total', 'cost', 'paid', 'teledyne', 'pack', 'two', 'piks', 'place', 'four']

which, unfortunately, is longer than 14 words and is dropped.

¹The outputs are given by demo_Preprocessing part in ipynb file

Model: 2. Word Embedding Skip-Gram Model



X : A token in a specific sentence.

$E(Y)$: The expected probability of another within the 'window' of X

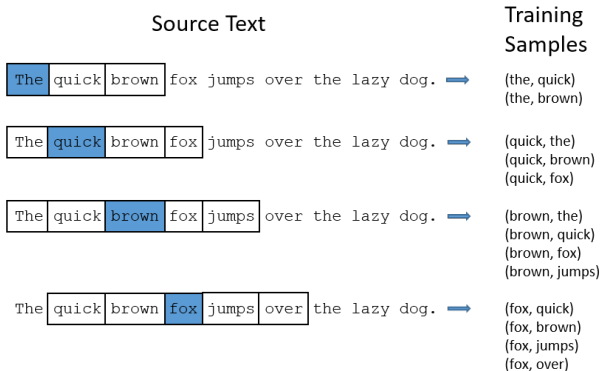


Figure: Figure from Word2Vec Tutorial

Model: 2. Word Embedding Skip-Gram Model




Figure: Figure from Wevi with the specific example. In *gesim*, the default return is the input vector.

Model: 2. Word Embedding Skip-Gram Model



Let I is the number of words in the corpus, V is the set of all feasible words in the vocabulary, and v_i are d -dimensional real-vector word embeddings. c be the window size for the estimation. We select the v_i to maximize:

$$\frac{1}{I} \sum_{i=1}^I \sum_{-c \leq j \leq c, j \neq 0} \log p(\text{word}_{i+j} \mid \text{word}_i)$$

where:

$$p(\text{word}_j \mid \text{word}_i) = \frac{\exp(v_j \cdot v'_i)}{\sum_{k=1}^{|V|} \exp(v_k \cdot v'_i)}.$$

Now we have the vector for each **token**. We concatenate the tokens at the sentence level to obtain the vector representation for a **sentence**:

$$v = [v_1, \dots, v_n] \in \mathbb{R}^{d \times n}$$

Model: 2. Word Embedding Skip-Gram Model



After the cleaning and tokenizing, we have²:

- ▶ ['well', 'really', 'know', 'well', 'pack', 'teledyne', 'piks', 'brj', 'works', 'even_though', 'clearly', 'ordered', 'goodman', 'without', 'notice', 'sent', 'official', 'waterpik', 'pack', 'jt', 'e', 'place', 'total', 'cost', 'paid', 'teledyne', 'pack', 'two', 'piks', 'place', 'four']

Finally, the Word2Vec turns it into a 20 by 1 vector (rounded by 4 digits, averaged over all the tokens for simplicity):

- ▶ [0.3674, -0.1955, 0.3926, -0.0273, -0.2642, 0.2046, 0.0489, 0.7512, -0.7463, 0.5247, 0.0127, -0.0863, 0.3502, -0.2463, 0.3582, 0.2546, 0.6332, -0.1017, -0.3856, -0.6695]

²The outputs are given by demo_Preprocessing part in ipynb file

Model 3: CNN Architecture Cont.



Figure 2. Convolutional Neural Network Architecture for Sentence Classification




Figure: Figure 2 from Timoshenko and Hauser (2019)

- ▶ The word embedding is not hard-coded, but with a `nn.embedding()`
- ▶ Padding was done at the end. I use default torch setting to pad on both sides.

For more tech details in the code see the replication files.

Model 3: Identify Informative Content with CNN



Numerical Representations of Words:

- ▶ Words are represented as real-valued vectors v_i from pre-trained embeddings.
- ▶ Sentence embedding:

$$v = [v_1, \dots, v_n] \in \mathbb{R}^{d \times n},$$

Convolution Layer:

- ▶ Applies filters $w_t \in \mathbb{R}^{d \times h_t}$ of size h_t to generate feature maps c^t .
- ▶ Feature computation:

$$c_i^t = \sigma(w_t \cdot v_{i:i+h_t-1} + b_t),$$

where $v_{i:i+h_t-1} = [v_i, \dots, v_{i+h_t-1}]$ and $\sigma(x) = \max(0, x)$ is ReLU.

Model 3: CNN Architecture



Pooling Layer:

- Performs global max pooling to summarize feature maps:

$$z_t = \max(c_1^t, \dots, c_{n-h_t+1}^t),$$

resulting in $z = [z_1, \dots, z_r]$, where r is the total number of filters across all filter sizes.

Softmax Layer:

- Outputs probabilities for sentence classification as informative or not:

$$p(y|z) = \text{Softmax}(z \cdot w + b).$$



A **direct replication** is hard in this scenario:

- ▶ The original code sent by Prof. Timoshenko to me in email was written in Python 2
- ▶ Based on an old version of keras for CNN
- ▶ `np.random.seed(100)` is defined in a loop.

So I re-write the codes in the replication file

- ▶ Switched the framework to the latest PyTorch under Python 3.10
- ▶ Replaced the loop with parallel computing for figure 5 and 7.
- ▶ Replaced old stats packages (e.g. fastcluster) with scipy
- ▶ For simplicity, I only replicate the **ML-related** sections in the main paper. not the non-ML tables and Appendix.

Model 4: Sample Diverse Content (with Replication)



Idea:

- ▶ We want to reduce the redundant sentences BEFORE e.g. manual consumer need analysis
- ▶ The following 3 sentences are considered redundant.
 - "When I am done, my teeth do feel'squeaky clean."*
 - "Every time I use the product, my teeth and gums feel professionally cleaned."*
 - "I am still shocked at how clean my teeth feel."*

Use sentence embeddings to reduce redundancy


- ▶ sampling content for manual review from maximally different parts of the space of sentence embeddings.

Model 4: Sample Diverse Content (with Replication)



Use the embedding from the Skip-Gram model, apply PCA to 2D by plotting the first 2 principle components.

Figure 6. Projections of 20-Dimensional Embeddings of Sentences onto Two Dimensions (PCA)




*** Shopping/Product Choice ●●● Strong Teeth and Gums

Note. Dots and crosses indicate analyst-coded primary customer needs.

(a) Figure 6

2D Projection of Sentence Embeddings (Classes 2 vs. 6)



■ Shopping/Product Choice ● Strong Teeth and Gums

(b) A Replication by Shuyi

Model 5 and so: Results and Replications



How good is the CNN model:


- ▶ Given different training sample?
- ▶ Compared with different models?

Model 5 Fig 5: Results and Replications




How good is the CNN model?

Figure 5. F1-score as a Function of the Size of the Training Sample



(a) Figure 5



(b) A Replication by Shuyi

"Performance of the CNN stabilizes after 500 training sentences, with some slight improvement after 500 training sentences."

Model 5 Table 2: Results and Replications Cont



How good is the CNN model?

Table 2. Alternative Machine-Learning Methods to Identify Informative Sentences

Method	Precision (%)	Recall (%)	Accuracy (%)	F ₁ (%)
Convolutional neural network (CNN)	74.4	73.6	74.2	74.0
CNN with asymmetric costs ($\gamma = 3$)	65.2	85.3	70.0	74.0
Recurrent neural network-LSTM	72.8	74.0	73.2	73.4
Multichannel CNN	70.5	74.9	71.8	72.6
Support vector machine	63.7	67.9	64.6	65.7

(a) Table 2

```
Results:
Model      Precision    Recall      F1          Accuracy
CNN (gamma = 1)  0.7500      0.7104      0.7297      0.7332
CNN (gamma = 3)  0.6548      0.8933      0.7557      0.7073
LSTM Network    0.7682      0.7301      0.7487      0.7516
SVM             0.7151      0.7296      0.7223      0.7157
(ds_torch) (base) liushijian@10-21-196-38 code %
```

(b) A Replication by Shuyi

- Results align with the original paper.
- Fail to run the Kim et al. 2014 Multichannel CNN as the tensor shape is tricky to align.

Model 5 Figure 7: Results and Replications Cont



Compare content selection approaches in terms of the **expected number of unique customer needs identified** in Y sentences.

- ▶ Unique Customer Needs: identified_needs.csv
- ▶ (1) randomly split the 6,700 preprocessed sentences, which are neither too short nor too long, into 3,700 training and 3,000 holdout samples
- ▶ 2) train the CNN using the training sample
- ▶ 3) draw Y sentences from the holdout sample for review. We count the unique needs identified in the Y sentences and repeat the process 10,000 times.

Three methods:


- ▶ Benchmark: Do nothing. Sample and count.
- ▶ Info Sentences CNN: Use CNN to identify informative sentences; sample from informative sentences for review.
- ▶ Redundancy Reduction: Sentence-embedding clusters to reduce redundancy after Info CNN ³

³In the code the authors then sample from each cluster (in total max_clusters generated by Unsupervised learning)

Model 5 Figure 7: Results and Replications Cont




Figure 7. Efficiencies Among Various Methods to Select UGC Sentences for Review



(a) Figure 7

Efficiencies Among Various Methods to Select UGC Sentences for Review



(b) A Replication by Shuyi

"The CNN improves efficiency as indicated by the dotted line. Using the CNN and clustering sentence embeddings increases efficiency further."



Empirical:

- ▶ Challenges the traditional Voice-of-Customer (VOC) methods. UGC can be an equally or more effective source for identifying customer needs.
- ▶ Empirical evidence: UGC + ML + Analyst captures a broader and more diverse set of customer needs compared to traditional methods.
- ▶ Very **low cost**, very effective, with very **simple structure**.

Methodology:

- ▶ Introduces a CNN-based classifier to filter out non-informative content from UGC.
- ▶ Utilizes sentence embeddings to cluster and reduce redundancy.
- ▶ "Human-AI Interaction"?

Limitations



Methodology:

- ▶ Explicitly perform word embedding and the CNN architecture are somehow "out of date" Vaswani (2017)
 - Restrictions on input length (Very long customer reviews can be useful)
 - Position encoding & End-to-end embedding.
 - Unnecessary to remove stop words and numbers beforehand.
- ▶ May not work for consumer needs search for innovative and new products. (VOC?)
 - Only works for firms with large UGC e.g. huge oral care product reviews
 - The economics of scale v.s. VOC.
- ▶ Do we really need to train a model for ourselves? Or can we solve it with LLM?

Extensions:


- ▶ Other sources of UGC: e.g. Stream Videos?
- ▶ Combined analysis: What are the "most urgent" consumer needs that e.g. result in frequent returns?

Why not just LLM?

Customer reviews

★★★★☆ 4.4 out of 5

19,870 global ratings



[How customer reviews and ratings work](#)

Review this product

Share your thoughts with other customers

[Write a customer review](#)

Cuisinart

Boil water in an instant



Customers say

Customers find the electric food mill grinder easy to use and a good value. It grinds coffee beans and nuts into fine powder. They appreciate its simple operation and user-friendly features. However, some customers have differing opinions on its durability, ease of cleaning, and design.

AI-generated from the text of customer reviews

Select to learn more

✓ Functionality | ✓ Grind quality | ✓ Value for money | ✓ Ease of use | ✓ Noise level |
[Durability](#) | [Ease of cleaning](#) | [Design](#)

Reviews with images

[See all photos](#)



[Top reviews](#)

Top reviews from the United States



Brandi Walsh

★★★★★ Great!

Reviewed in the United States on January 28, 2025

Style: Modern | Color: Stainless Steel | [Verified Purchase](#)

Figure: LLM Summary Example from Amazon

New Methods: Why not just ChatGPT



Imagine you are a marketing expert. Now I am giving you two datasets with user generated contents on the oral products.. The "Sentence Text" are the user generated reviews, the "isNeed" is a dummy variable that equals 1 if the "Sentence Text" is talking about a specific need that the consumer wants. Now, please read the "train_4000.csv" carefully. Please fully understand what does it mean for an review to represent consumer need. Do you understand the task?

Great. Now I will give you another file. Please read this file carefully, and add a new column called isNeed_res that equals 1 if a sentence is about a specific need, and 0 otherwise. Please use the knowledge you obtained from the previous file. Please return a csv file with only two columns: "Sentence ID" and "is-Need_res".

Train/Test = 1

Accuracy: 0.48325

Probably not that easy.



- Timoshenko, A., & Hauser, J. R. (2019, January). Identifying Customer Needs from User-Generated Content. *Marketing Science*, 38(1), 1–20. Retrieved 2025-01-31, from <https://pubsonline.informs.org/doi/10.1287/mksc.2018.1123> doi: 10.1287/mksc.2018.1123
- Vaswani, A. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*.